

Hierarchical Bayes models for daily rainfall time series at multiple locations from different data sources

1 Motivation

Crop yields respond to rainfall and temperature, but average decadal or even average monthly rainfall can obscure the effects of weather on yields. Frequent exposure to extreme temperatures (high and low) matters tremendously (Schlenker 2009). Furthermore, the daily variation in other variables, such as soil moisture, has been found to be equally important, and if not accounted for, allows for temperature's effect to be overstated (Ortiz 2014). Specifically, the NUMBER of times and WHEN a crop is exposed to highs/lows in temperature and moisture is crucial, which decadal or monthly measures cannot capture. [Is this disingenuous if we don't have temperature but just rainfall amount?]

Daily statistics and extremes of weather often have a greater impact on end-users than climatological averages. For example, although seasonal rainfall totals and average temperatures can provide a general guide to crop yields, plants are also extremely sensitive to extreme weather events over a few hours (Schlenker, 2009, Ortiz 2014), especially during flowering or pod-set. A classic example is in groundnut, where high temperatures during the one or two days of flowering leads to pod death and plant death (Wheeler et al, 1996; Wheeler et al, 2000), or in maize, where a dry spell over the short Anthesis Tasselling Interval has a well measured and catastrophic impact on yields (Bolaos and Edmeades (1996)). Similar impacts can be seen in hydrology, where extreme rainfall events that are often of more interest to flood modellers than seasonal averages (REFERENCE), or in climate science, where it is crucial to be able to convert projections of regional averages into changes in local weather (REFERENCE).

The ability to generate realistic daily time-series of weather is therefore valued in many end-user applications such as hydrological, climate or crop simulation modelling, especially over regions of the world where there are sparse ground based weather station networks. The technique for producing such time series has moved from deterministic weather models where rainfall follows a stochastic markov process where transition probabilities dictate the chance of rain, give only the particular state. As in (weather generation game) the seminal papers in this area are xxxx. The paper also reviews the techniques of XXX and YYYY[HELEN]. The objective in any of these approaches is to be able to simulate long synthetic time-series of weather that reflect key observed statistics in the region of interest e.g. means, dry spell lengths, or probability of extreme events occurring.

The Hierarchal Bayes method goes beyond the markov approach for several reasons: [assumes simplistic relationships between weather variables, and regional climate modelling is still in its infancy over many regions, with well documented problems in modelling variables such as rainfall. But you guys can incorporate all the information you have, so multiple time-series at a location, SSTs etc without knowing the exact prior relationship]

Recently, there has been a high level of interest in spatially correlated weather generators, which would have a clear benefit when considering modelling river basins or regional food security. Incorporating spatial correlation is also key when one considers that most weather

events such as droughts or floods have a high impact precisely because they are spatially correlated, regional events, thus incorporating observed spatial correlations into a well calibrated stochastic approach is likely to lead to more realistic impact models. Several different techniques and methods have been suggested. The MarkSim weather generator uses spatially correlated input grids of statistics to allow one to allow estimation at a site where there are no observations (Jones and Thornton, 1993). Wilks (1999,1998) extended a richardson weather generator by driving it with a grid of spatially correlated random numbers. More recently this approach was extended in the GIST weather generator, incorporates spatial correlation through the use of correlation matrices to sample from a cumulative probability function at each location (Baigorria and Jones, 2010), while Kleiber et al (2012) suggests the use of latent Gaussian processes. In 2012. Greatrex (2012) proposed a geostatistical sequential simulation approach coupled to a markov generator, which would allow spatially correlated ensembles of maps of rainfall to be generated. It is important to note that many of these approaches rely on a large amount of observed and calibrated data, for example a dense station network to create variograms [EXPLAIN].

Then move onto the heirarchical bayes approach and discuss the sanso guenni and related papers.....

The hierchacal bayesian estimation is a modeling approach first employed in XXXXX (KENNY?). The approach has two steps or levels: the parameters describing the effect of say satellite information in describing rainfall, and then given those parameters, the probability that rainfall actually occurred. We also assume that given the parameters, the probability of rainfall follows a particular model such as a probit. Use initial crude estimates of betas, an interative process called Gibbs sampling is used to estimate the parameter conditional on the current estimates of the other parameters by making a random draw from the conditional distribution. After a burn in period, the estimates from several different "chains" converge to the correct estimates for each parameter.

The approach has been broadened to incorporate features such as non-stationarity sanso kleiber

The novelty of our approach is our ability to account for two levels of the rainfall time series information at each site: 1) both the location of the measurement and 2) the instrument used to measure, as well to incorporate. In addition, we 3) estimate the the spatial covariance between all these series and 4) the noise or error in recording rainfall due to the instrument itself. By incorporating several layers of information sources for each site in which we would like to predict weather, we are adding more information to the parameter estimates and, hopefully, improving out-sample predictions. Because we used daily data to fit our model, our simulated data and out-of-sample predictions are daily, which allows us to compute statistics sensitive to daily measurements: probability of rainfall, dry spells, and extreme rainfall.

The remainder the paper is as follows. In Section

2 Procedure

We model a set of 15 time series of daily rainfall in Ethiopia during the time period from 1992 to 2010, where these 15 time series come from six different locations, and each location has at least two time series of daily rainfall associated with it. The reason we observed multiple time series for each location is that for each location we have multiple sources of data, including rain station data and satellite-based rainfall proxy data. The statistical challenge to modeling such data is to separately estimate the spatial variability between locations and the within-location measurement-based variability. A standard hierarchical model with sets of 15 exchangeable parameters – one for each time series – would conflate these two sources of variation. The model we introduce here – a hierarchical Bayesian model with one level for locations, and another level for multiple data sources within a location – explicitly models each of these two sources of variation.

Figure 1 shows the six locations from which the daily rainfall time series are measured. The names of the locations are Hagere Selam, Maykental, Mekele, Abi Adi, Agibe, and Adi Ha. The last of these, Adi Ha, is the location we are most interested in, because we wish to provide rainfall insurance for farmers who live there. Specifically, we want to model rainfall at one of the automated rain stations at Adi Ha, which is one of the 15 time series in our data set, but is also the series that has the least amount of observed data – only about 200 days worth of data from 2009.




Figure 1: A map of the 6 locations, where the green squares denote ARC pixels and the pins denote rain station locations. The inset in the upper right corner shows the area of the map with a red rectangular box; this region is in north central Ethiopia.

The specifics of the data are as follows:

- For the first five locations, we have rain station data as well daily measurements from a satellite product called ARC, which is a rainfall proxy based on the temperature of the clouds over an area of about one hundred square kilometers. This comprises $5 \times 2 = 10$ time series.
- For the sixth location, Adi Ha, we have five separate data sources:
 1. One reliable rain station from which we only have 200 days of data from 2009-2010; this is the time series in which we have the most interest, because it is a new, accurate rain station on which we want to base insurance contracts.
 2. One unreliable rain station from which we have about 7 years of data from 2000-2009, with about 2 years of missing data interspersed.
 3. The ARC satellite proxy.
 4. Two additional satellite proxies that are different from ARC.

Figure 2 shows the range of observed and missing data for each of these 15 time series; note that the time scale goes back to 1961 for one of the rain stations, but for simplicity, we only consider the time span of 1992-2010 in our model fit, because this span contains most of the data.




Figure 2: A visualization of the observed data for each of the 15 time series we model. The black hash marks denote rain station data, and the red hash marks denote satellite-based data.

Table 1 contains some background information and summary statistics related to each time series of daily rainfall. For each time series we record the latitude, longitude, and elevation of the location where measurements were made, and the number of days of observed data. The maximum distance between locations is about 70 kilometers (between Mekele in the southeast and Maykental in the northwest).

Table 1: Background information about the 15 time series

	Site	Latitude	Longitude	Elev. (m)	Num. Obs
1	Hagere Salaam	13° 38' 49"	39° 10' 19"	2625	4887
2	Hagere Salaam (ARC)	"	"	"	5632
3	Maykental	13° 56' 13"	38° 59' 49"	1819	5620
4	Maykental (ARC)	"	"	"	5632
5	Mekele	13° 28' 1"	39° 31' 1"	2247	6205
6	Mekele (ARC)	"	"	"	5632
7	Abi Adi	13° 37' 19"	39° 0' 10"	1849	4205
8	Abi Adi (ARC)	"	"	"	5632
9	Agibe	13° 33' 43"	39° 3' 43"	1952	4722
10	Agibe (ARC)	"	"	"	5632
11	Adi Ha (ARC)	13° 43' 48"	39° 05' 38"	1713	5632
12	Adi Ha (Rain Station - Manual)	"	"	"	2769
13	Adi Ha (RFE2)	"	"	"	2920
14	Adi Ha (CMorph)	"	"	"	2190
15	Adi Ha (Rain Station - Automatic)	"	"	"	186

3 Exploratory Data Analysis

In this part of Ethiopia, the rainy season lasts roughly from June to October. Figure 3 shows the percentage of rainy days and the average amount of rain as a function of the time of year for each time series. The basic modeling strategy will be to use a set of periodic functions to model rainfall as a function of the season of the year.

We are also interested in the difference, on average, between the measurements of rainfall based on the ARC satellite proxy and the rain stations. Comparing rainfall frequencies pooled over 5-day periods, averaging across all parts of the year and all five locations with exactly one rain station and one ARC measurement, we find that the ARC records about 3% fewer days of rainfall than the rain stations. Across locations, this difference ranges from about -6% (Hager Selam) to +1% (Agibe).

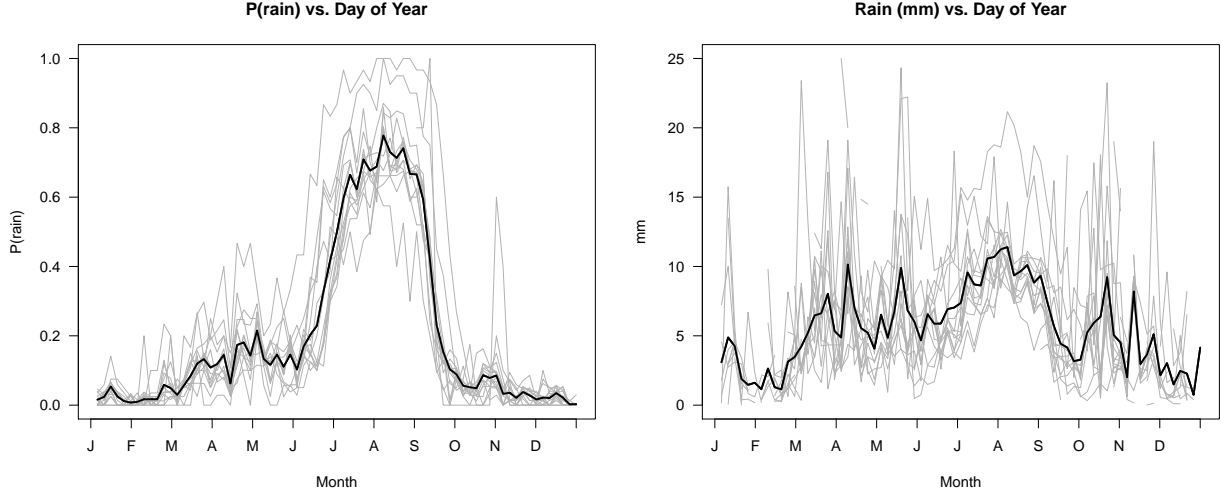


Figure 3: Plots of the percentage of rainy days (pooled into 5-day bins) and the average amount of rain as a function of the season. The gray lines are for each of the 15 individual time series, and the black lines are averaged across all 15 time series.

4 The Model

We fit a tobit model for daily rainfall at multiple locations, with multiple time series observed at each site. Let us first set up some notation. Let $S = 6$ denote the number of locations where we measure rainfall, and $\mathbf{J} = \{2, 2, 2, 2, 2, 5\}$ is the vector denoting the number of daily rainfall time series observed for each of the S locations. The total number of days in our time series is $N = 6679$ days, from 1/1/1992 through 7/28/2010. Let Y_{stj} denote the amount of rainfall, measured in mm, for location $s \in (1, \dots, S)$, day $t \in (1, \dots, N)$, and time series $j = (1, \dots, J_s)$. Last, let D_{ik} denote the Euclidian distance between site i and k , for $i, k \in (1, \dots, S)$.

We model Y_{stj} using a hierarchical Bayesian tobit regression model, where the levels of the hierarchy correspond to different sources of variation:

$$Y_{stj} = \begin{cases} W_{stj} & \text{if } W_{stj} > 0 \\ 0 & \text{if } W_{stj} \leq 0, \end{cases} \quad \text{Observed rainfall}$$

$$\mathbf{W}_{st} \sim N_{J_s}(\mathbf{Z}_{st} + \mathbf{X}_s^{\text{ARC}} \beta_s^{\text{ARC}}, \frac{1}{\gamma_{st}} \Sigma_s), \quad \text{Latent rainfall}$$

$$\mathbf{Z}_t \sim N_S(\mathbf{X}_t \beta^Z, \tau^2 \mathbf{V}), \quad \text{Spatial mean rainfall}$$

$$\beta_{ps}^Z \sim N(\mu_p, \sigma_p^2),$$

$$\mu_p \sim N(0, 5^2),$$

$$\sigma_p \sim \frac{1}{2}t(0, 1, \text{df} = 3),$$

$$\tau \sim \frac{1}{2}t(0, 1, \text{df} = 3),$$

$$\mathbf{V} = \{v_{ik}\}, v_{ii} = 1, v_{ik} = \exp(-\lambda d_{ik}) \quad \text{for } i, k \in (1, \dots, S)$$

$$\lambda \sim \text{gamma}(\text{shape} = 50, \text{scale} = 3),$$

$$\mathbf{X}_t = (1, t, t^2, \sin(2\pi t\omega_1), \cos(2\pi t\omega_1), \dots, \sin(2\pi t\omega_4), \cos(2\pi t\omega_4), \\ X_{\text{Jan}_{[t]}}^{\text{nino}}, X_{\text{Feb}_{[t]}}^{\text{nino}}, \dots, X_{\text{Dec}_{[t]}}^{\text{nino}}),$$

$$\mathbf{X}_{sj}^{\text{ARC}} = 1(\text{jth time series at site } s \text{ is an ARC product}),$$

$$\beta_s^{\text{ARC}} \sim N(\mu^{\text{ARC}}, \tau_{\text{ARC}}^2),$$

$$\mu^{\text{ARC}} \sim N(0, 5^2),$$

$$\tau_{\text{ARC}} \sim \frac{1}{2}t(0, 1, \text{df} = 3),$$

$$\Sigma_s \sim \text{Inv-Wish}(v_0 = J_s, \Lambda_0^{-1} = \text{diag}(J_s))$$

$$\gamma_{st} \sim \text{gamma}(\text{shape} = \frac{5}{2}, \text{scale} = \frac{2}{5}),$$

The explanation of the model is as follows.

1. The first level of the model is a standard tobit regression, where we model the observed rainfall, Y_{stj} , as being equal to the j^{th} component of the latent rainfall vector, \mathbf{W}_{st} , if it is greater than zero, and equal to zero if the j^{th} component of the latent rainfall vector is less than or equal to zero.
2. Next, for each location s , the length- J_s vector of latent rainfall amounts, \mathbf{W}_{st} , is a multivariate t random variable centered at the spatial mean rainfall amount for that location, \mathbf{Z}_{st} , and offset by an ARC bias effect, β_s^{ARC} (where X_{sj}^{ARC} is an indicator variable of whether time series j at location s is an ARC satellite product). The latent rainfall, \mathbf{W}_{st} , is a multivariate- t random variable because it is a scale mixture of multivariate normals with a mixture weight, $\frac{1}{\gamma_{st}}$, for the covariance, Σ_s , that is drawn from a gamma distribution.
3. The location-specific covariance matrices Σ_s allow the multiple time series at each location to be correlated in unique ways. The mixing parameters γ_{st} determine the widths of the tails of the multivariate- t distributions, and are modeled with a gamma prior distribution with shape and scale parameters equal to $5/2$ and $2/5$, respectively, which implies that the multivariate- t distribution, \mathbf{W}_{st} , has 5 degrees of freedom. (In

follow-up models, we could relax this assumption and estimate from the data how heavy the tails should be; the choice of 5 degrees of freedom is based on the fits of some simple, exploratory models).

4. The spatial mean rainfall amount for day t , \mathbf{Z}_t , is modeled as a multivariate normal random variable whose mean depends on the day, t (linearly, quadratically, and periodically, with periods $\omega = \frac{1}{365}(1, 2, 3, 4)$), and also on effects from El Nino, where the El Nino effect is an additive constant that depends on the month (allowing El Nino to have different effects across the 12 months of the year).
5. The covariance matrix of \mathbf{Z}_t , $\tau^2 \mathbf{V}$, depends on τ , a scaling factor, one known input, the Euclidean distance between locations, and one unknown parameter, λ . The spatial correlation in rainfall between locations is modeled separately from the noise inherent in the different measurement methods at each location, which is modeled by Σ_s . The model assumes that the covariances of pairs of Z_{st} 's decay exponentially with the Euclidean distance between the pairs of locations at the unknown rate λ , which we estimate from the data using a relatively flat prior.
6. The rest of the model is straightforward. We shrink the β_{ps}^Z 's for each location toward a common mean μ_p . We also model the ARC biases, β_s^{ARC} as normal random variables with an unknown mean, μ_{ARC} , and variance τ_{ARC}^2 .

5 Fitting the Model

We used MCMC to fit the model. There was a combination of Gibbs steps and metropolis steps. It took about [x] minutes to sample 3 chains of 5000 iterations each. Details of the MCMC implementation are available in the appendix.

6 Simulation Study

First, we ran a simulation study, where we simulated data from a set of known parameters. These parameter values were chosen to approximate values that could have produced our observed data, based on EDA. The size of the simulated data set was similar to the real data in both dimensions (the number of time series and the number of days).

The results were very encouraging: we were able to precisely estimate the spatial correlation between locations as well as the variability of different data sources within single locations. Figure 4 shows trace plots for λ , μ_α , and τ_α . All three parameters were well-estimated, and convergence was relatively quick. Furthermore, we were able to accurately estimate different ARC biases, α_s for each location, as well as different amounts of variability at each location, τ_s . Last, the estimates of β were accurate.




Figure 4: Simulation study results for λ , μ_α , and τ_α , where the horizontal blue lines represent the known true values of the parameters.

7 Results

The results on the real data are equally encouraging.

1. prob and mean rainfall historical vs. sim
2. dry spell historical vs sim
3. rainfall above a certain amount
3. metrics showing prob of rainfall in all locations?

8 Out-of Sample Predictions

We also generate out of sample rainfall data using 2011's el nino data, and sampling from the posterior distribution from our initial run. We can compare these data to the actual weather time series for 2011 for: